

Breadth First Search, Dijkstra's Algorithm for Shortest Paths

Lecture 17

October 30, 2018

Part I

Breadth First Search

Breadth First Search (BFS)

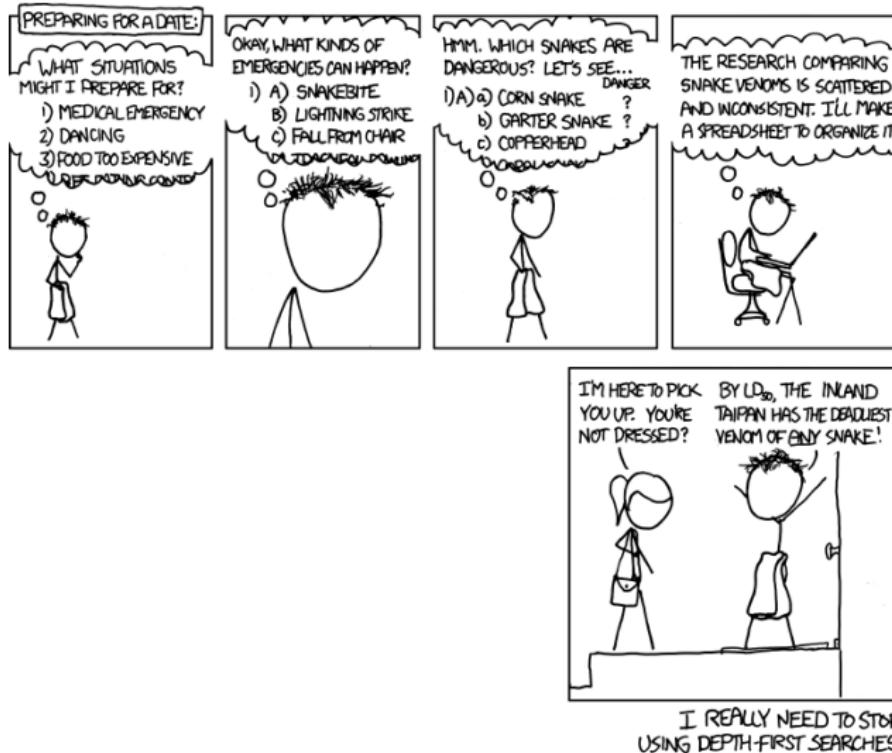
Overview

- (A) **BFS** is obtained from **BasicSearch** by processing edges using a data structure called a **queue**.
- (B) It processes the vertices in the graph in the order of their shortest distance from the vertex **s** (the start vertex).

As such...

- ① **DFS** good for exploring graph structure
- ② **BFS** good for exploring *distances*

xkcd take on DFS



Queue Data Structure

Queues

A **queue** is a list of elements which supports the operations:

- ① **enqueue**: Adds an element to the end of the list
- ② **dequeue**: Removes an element from the front of the list

Elements are extracted in **first-in first-out (FIFO)** order, i.e., elements are picked in the order in which they were inserted.

BFS Algorithm

Given (undirected or directed) graph $G = (V, E)$ and node $s \in V$

BFS(s)

Mark all vertices as unvisited

Initialize search tree T to be empty

Mark vertex s as visited

set Q to be the empty queue

enq(s)

while Q is nonempty **do**

$u = \text{deq}(Q)$

for each vertex $v \in \text{Adj}(u)$

if v is not visited **then**

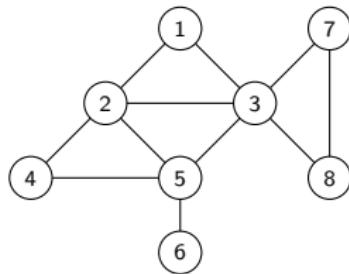
 add edge (u, v) to T

 Mark v as visited and **enq(v)**

Proposition

BFS(s) runs in $O(n + m)$ time.

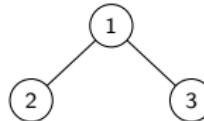
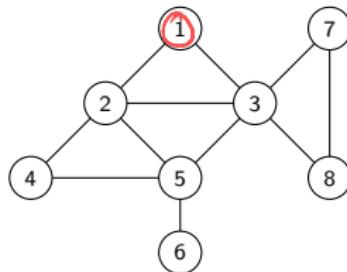
BFS: An Example in Undirected Graphs



1

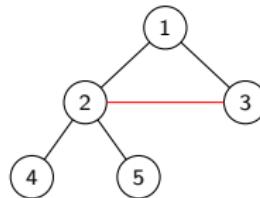
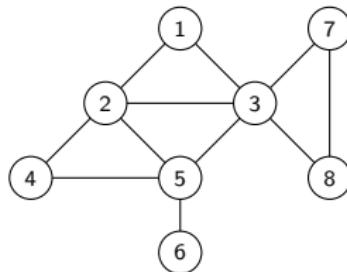
1. [1]

BFS: An Example in Undirected Graphs



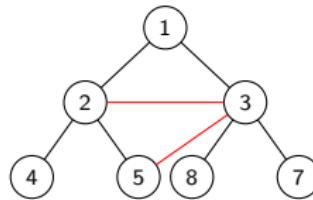
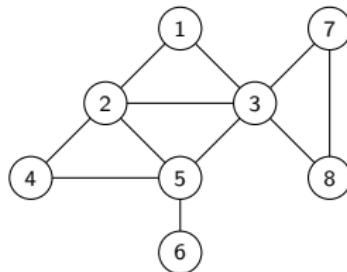
1. [1]
2. [2,3]

BFS: An Example in Undirected Graphs



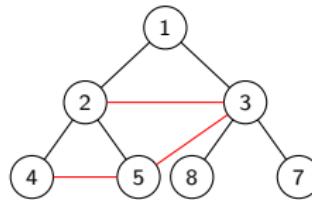
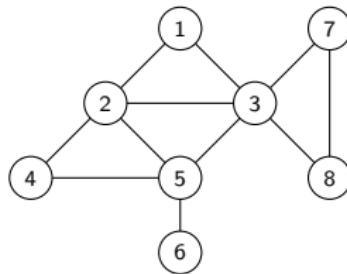
1. [1]
2. [2,3]
3. [3,4,5]

BFS: An Example in Undirected Graphs



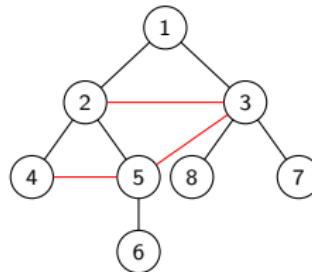
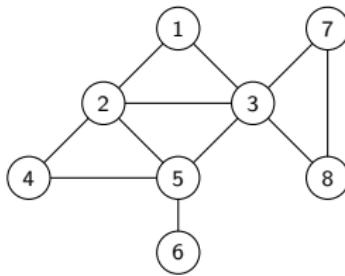
1. [1]
2. [2,3]
3. [3,4,5]
4. [4,5,7,8]

BFS: An Example in Undirected Graphs



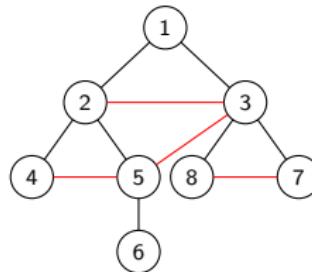
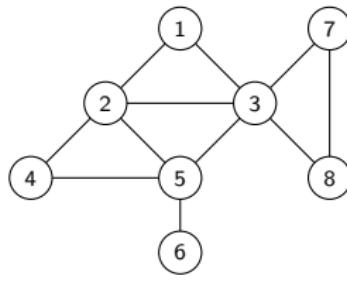
1. [1]
2. [2,3]
3. [3,4,5]
4. [4,5,7,8]
5. [5,7,8]

BFS: An Example in Undirected Graphs



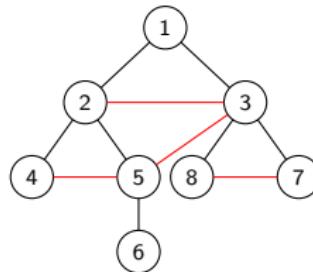
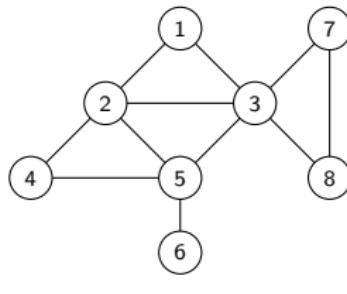
1. [1]
2. [2,3]
3. [3,4,5]
4. [4,5,7,8]
5. [5,7,8]
6. [7,8,6]

BFS: An Example in Undirected Graphs



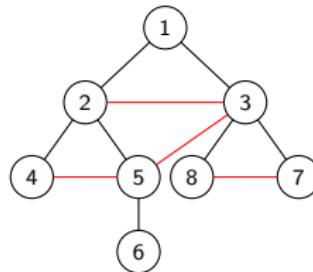
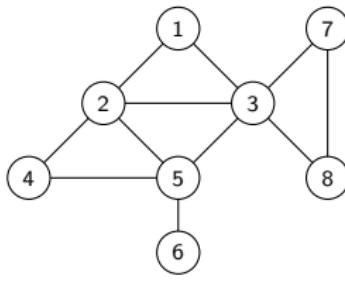
1. [1]
2. [2,3]
3. [3,4,5]
4. [4,5,7,8]
5. [5,7,8]
6. [7,8,6]
7. [8,6]

BFS: An Example in Undirected Graphs



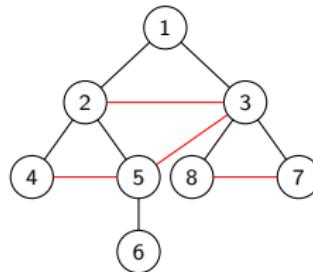
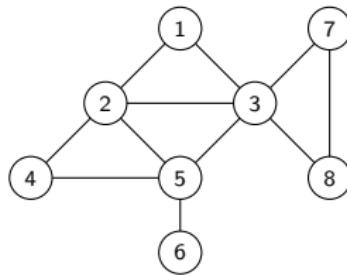
- | | | |
|------------|--------------|----------|
| 1. [1] | 4. [4,5,7,8] | 7. [8,6] |
| 2. [2,3] | 5. [5,7,8] | 8. [6] |
| 3. [3,4,5] | 6. [7,8,6] | |

BFS: An Example in Undirected Graphs



1. [1]
2. [2,3]
3. [3,4,5]
4. [4,5,7,8]
5. [5,7,8]
6. [7,8,6]
7. [8,6]
8. [6]
9. []

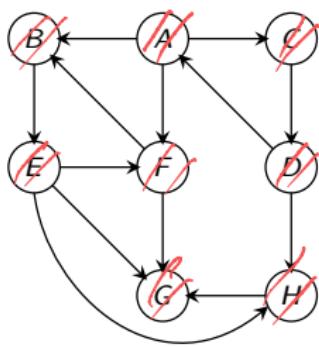
BFS: An Example in Undirected Graphs



- | | | |
|------------|--------------|----------|
| 1. [1] | 4. [4,5,7,8] | 7. [8,6] |
| 2. [2,3] | 5. [5,7,8] | 8. [6] |
| 3. [3,4,5] | 6. [7,8,6] | 9. [] |

BFS tree is the set of black edges.

BFS: An Example in Directed Graphs



$[A] \rightarrow \text{pop} \rightarrow []$

$[BCF] \rightarrow \text{pop} \rightarrow [CF]$

$[CFG] \rightarrow \text{pop} \rightarrow [FG]$

$[FEG] \rightarrow \text{pop} \rightarrow [EP]$

$[EDG] \rightarrow [DGH]$

$[GH]$ $\rightarrow [H]$

$H \rightarrow []$

BFS with Distance

BFS(s)

Mark all vertices as unvisited; for each v set $\text{dist}(v) = \infty$

Initialize search tree T to be empty

Mark vertex s as visited and set $\text{dist}(s) = 0$

set Q to be the empty queue

enq(s)

while Q is nonempty **do**

$u = \text{deq}(Q)$

for each vertex $v \in \text{Adj}(u)$ **do**

if v is not visited **do**

add edge (u, v) to T

 Mark v as visited, **enq(v)**

and set $\text{dist}(v) = \text{dist}(u) + 1$

Properties of BFS: Undirected Graphs

Theorem

The following properties hold upon termination of $\text{BFS}(s)$

- (A) *The search tree contains exactly the set of vertices in the connected component of s .*
- (B) *If $\text{dist}(u) < \text{dist}(v)$ then u is visited before v .*
- (C) *For every vertex u , $\text{dist}(u)$ is the length of a shortest path (in terms of number of edges) from s to u .*
- (D) *If u, v are in connected component of s and $e = \{u, v\}$ is an edge of G , then $|\text{dist}(u) - \text{dist}(v)| \leq 1$.*

Properties of BFS: Directed Graphs

Theorem

The following properties hold upon termination of $\text{BFS}(s)$:

- (A) *The search tree contains exactly the set of vertices reachable from s*
- (B) *If $\text{dist}(u) < \text{dist}(v)$ then u is visited before v*
- (C) *For every vertex u , $\text{dist}(u)$ is indeed the length of shortest path from s to u*
- (D) *If u is reachable from s and $e = (u, v)$ is an edge of G , then $\text{dist}(v) - \text{dist}(u) \leq 1$.*
Not necessarily the case that $\text{dist}(u) - \text{dist}(v) \leq 1$.

BFS with Layers

BFSLayers(s):

Mark all vertices as unvisited and initialize T to be empty

Mark s as visited and set $L_0 = \{s\}$

$i = 0$

while L_i is not empty **do**

 initialize L_{i+1} to be an empty list

for each u in L_i **do**

for each edge $(u, v) \in \text{Adj}(u)$ **do**

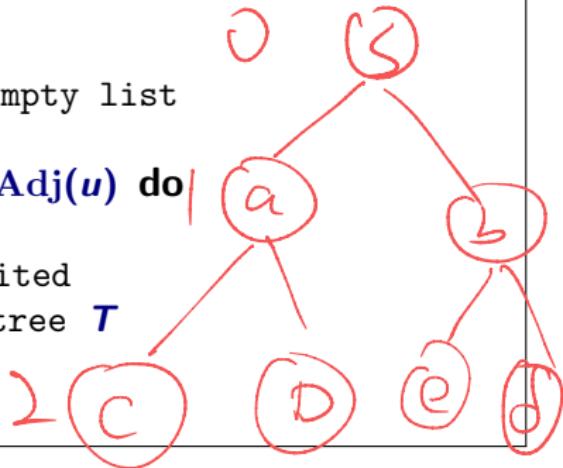
 if v is not visited

 mark v as visited

 add (u, v) to tree T

 add v to L_{i+1}

$i = i + 1$



BFS with Layers

BFSLayers(s):

Mark all vertices as unvisited and initialize T to be empty

Mark s as visited and set $L_0 = \{s\}$

$i = 0$

while L_i is not empty **do**

 initialize L_{i+1} to be an empty list

for each u in L_i **do**

for each edge $(u, v) \in \text{Adj}(u)$ **do**

 if v is not visited

 mark v as visited

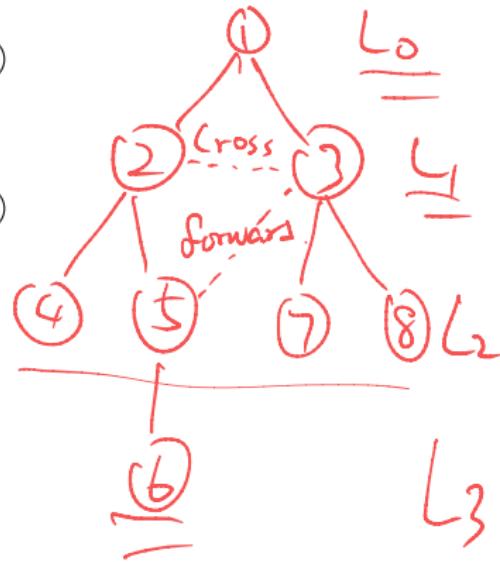
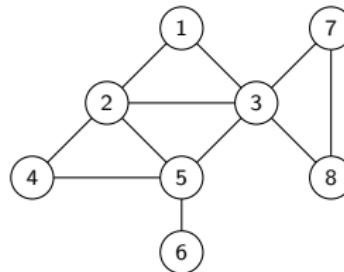
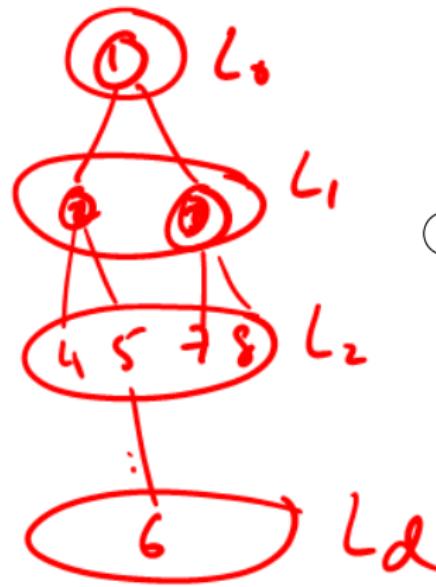
 add (u, v) to tree T

 add v to L_{i+1}

$i = i + 1$

Running time: $O(n + m)$

Example



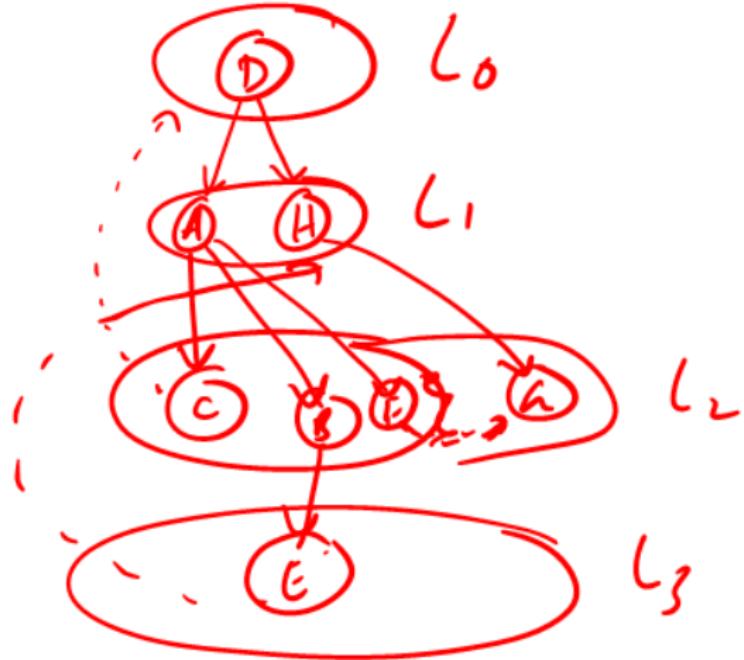
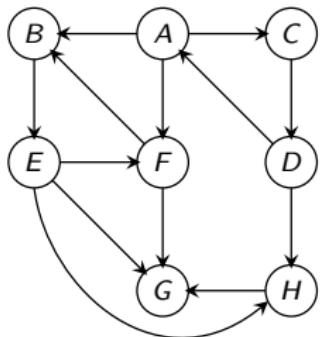
BFS with Layers: Properties

Proposition

The following properties hold on termination of $\text{BFSLayers}(s)$.

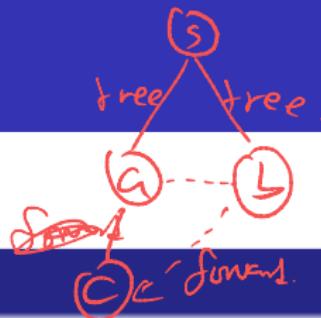
- ① $\text{BFSLayers}(s)$ outputs a **BFS** tree
- ② L_i is the set of vertices at distance exactly i from s
- ③ If G is undirected, each edge $e = \{u, v\}$ is one of three types:
 - ① **tree** edge between two consecutive layers
 - ② non-tree **forward/backward** edge between two consecutive layers
 - ③ non-tree **cross-edge** with both u, v in same layer
- ④ \implies Every edge in the graph is either between two vertices that are either (i) in the same layer, or (ii) in two consecutive layers.

Example



BFS with Layers: Properties

For directed graphs



Proposition

The following properties hold on termination of **BFSLayers**(s), if G is directed.

For each edge $e = (u, v)$ is one of four types:

- ① a **tree** edge between consecutive layers, $u \in L_i, v \in L_{i+1}$ for some $i \geq 0$
- ② a non-tree **forward** edge between consecutive layers
- ③ a non-tree **backward** edge
- ④ a **cross-edge** with both u, v in same layer

Part II

Shortest Paths and Dijkstra's Algorithm

Shortest Path Problems

Shortest Path Problems

Input A (undirected or directed) graph $G = (V, E)$ with edge lengths (or costs). For edge $e = (u, v)$, $\ell(e) = \ell(u, v)$ is its length.

- ① Given nodes s, t find shortest path from s to t .
- ② Given node s find shortest path from s to all other nodes.
- ③ Find shortest paths for all pairs of nodes.

Shortest Path Problems

Shortest Path Problems

Input A (undirected or directed) graph $G = (V, E)$ with edge lengths (or costs). For edge $e = (u, v)$, $\ell(e) = \ell(u, v)$ is its length.

- ① Given nodes s, t find shortest path from s to t .
- ② Given node s find shortest path from s to all other nodes.
- ③ Find shortest paths for all pairs of nodes.

Many applications!

Single-Source Shortest Paths:

Non-Negative Edge Lengths

Single-Source Shortest Path Problems

- ① **Input:** A (undirected or directed) graph $G = (V, E)$ with **non-negative** edge lengths. For edge $e = (u, v)$, $\ell(e) = \ell(u, v)$ is its length.
- ② Given nodes s, t find shortest path from s to t .
- ③ Given node s find shortest path from s to all other nodes.

Single-Source Shortest Paths:

Non-Negative Edge Lengths

Single-Source Shortest Path Problems

- ① **Input:** A (undirected or directed) graph $G = (V, E)$ with **non-negative** edge lengths. For edge $e = (u, v)$, $\ell(e) = \ell(u, v)$ is its length.
 - ② Given nodes s, t find shortest path from s to t .
 - ③ Given node s find shortest path from s to all other nodes.
-
- ① Restrict attention to directed graphs
 - ② Undirected graph problem can be reduced to directed graph problem - how?

Single-Source Shortest Paths:

Non-Negative Edge Lengths

Single-Source Shortest Path Problems

- ① **Input:** A (undirected or directed) graph $G = (V, E)$ with **non-negative** edge lengths. For edge $e = (u, v)$, $\ell(e) = \ell(u, v)$ is its length.
 - ② Given nodes s, t find shortest path from s to t .
 - ③ Given node s find shortest path from s to all other nodes.
-
- ① Restrict attention to directed graphs
 - ② Undirected graph problem can be reduced to directed graph problem - how?
 - ① Given undirected graph G , create a new directed graph G' by replacing each edge $\{u, v\}$ in G by (u, v) and (v, u) in G' .
 - ② set $\ell(u, v) = \ell(v, u) = \ell(\{u, v\})$
 - ③ Exercise: show reduction works. **Relies on non-negativity!**

Single-Source Shortest Paths via BFS

Special case: All edge lengths are **1**.

Single-Source Shortest Paths via BFS

Special case: All edge lengths are **1**.

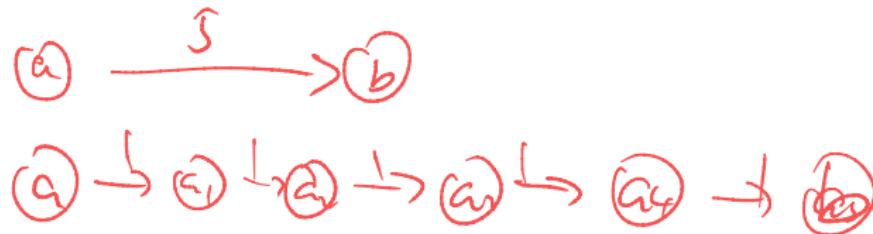
- ① Run **BFS(s)** to get shortest path distances from s to all other nodes.
- ② $O(m + n)$ time algorithm.

Single-Source Shortest Paths via BFS

Special case: All edge lengths are **1**.

- ① Run **BFS(s)** to get shortest path distances from s to all other nodes.
- ② $O(m + n)$ time algorithm.

Special case: Suppose $\ell(e)$ is an integer for all e ?
Can we use **BFS**?



Single-Source Shortest Paths via BFS

Special case: All edge lengths are 1.

- ① Run **BFS(s)** to get shortest path distances from s to all other nodes.
- ② $O(m + n)$ time algorithm.

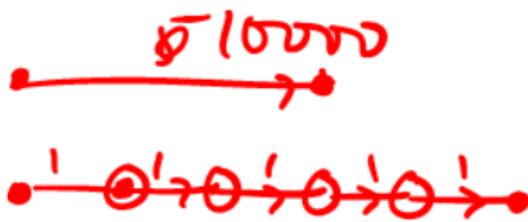
Special case: Suppose $\ell(e)$ is an integer for all e ?

Can we use **BFS**? Reduce to unit edge-length problem by placing $\ell(e) - 1$ dummy nodes on e

$$n, m$$

$$n' = n + \sum_e \ell(e)$$

$$n + m + \sum_e \lg(\ell(e))$$



Single-Source Shortest Paths via BFS

Special case: All edge lengths are 1 .

- ① Run **BFS**(s) to get shortest path distances from s to all other nodes.
- ② $O(m + n)$ time algorithm.

Special case: Suppose $\ell(e)$ is an integer for all e ?

Can we use **BFS**? Reduce to unit edge-length problem by placing $\ell(e) - 1$ dummy nodes on e

Let $L = \max_e \ell(e)$. New graph has $O(mL)$ edges and $O(mL + n)$ nodes. **BFS** takes $O(mL + n)$ time. Not efficient if L is large.

Towards an algorithm

Why does **BFS** work?

Towards an algorithm

Why does **BFS** work?

BFS(s) explores nodes in increasing distance from s

Towards an algorithm

Why does **BFS** work?

BFS(s) explores nodes in increasing distance from s

Lemma

Let G be a directed graph with non-negative edge lengths. Let $\text{dist}(s, v)$ denote the shortest path length from s to v . If $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$ is a shortest path from s to v_k then for $1 \leq i < k$:

- ① $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_i$ is a shortest path from s to v_i
- ② $\text{dist}(s, v_i) \leq \text{dist}(s, v_k)$. Relies on non-neg edge lengths.

Towards an algorithm

Lemma

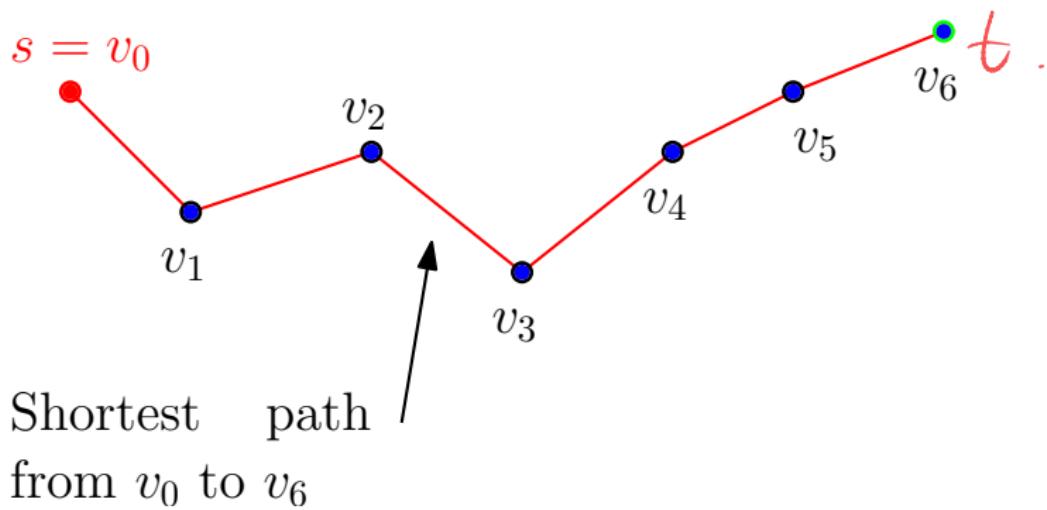
Let G be a directed graph with non-negative edge lengths. Let $\text{dist}(s, v)$ denote the shortest path length from s to v . If $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$ is a shortest path from s to v_k then for $1 \leq i < k$:

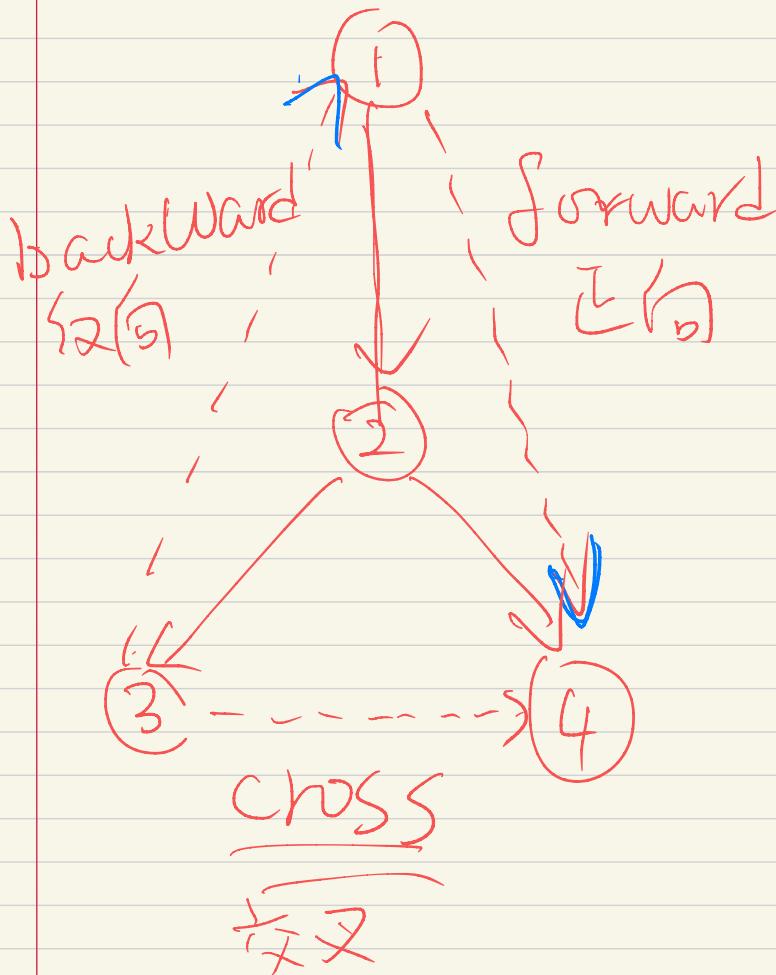
- ① $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_i$ is a shortest path from s to v_i
- ② $\text{dist}(s, v_i) \leq \text{dist}(s, v_k)$. *Relies on non-neg edge lengths.*

Proof.

Suppose not. Then for some $i < k$ there is a path P' from s to v_i of length strictly less than that of $s = v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_i$. Then P' concatenated with $v_i \rightarrow v_{i+1} \dots \rightarrow v_k$ contains a strictly shorter path to v_k than $s = v_0 \rightarrow v_1 \dots \rightarrow v_k$. For the second

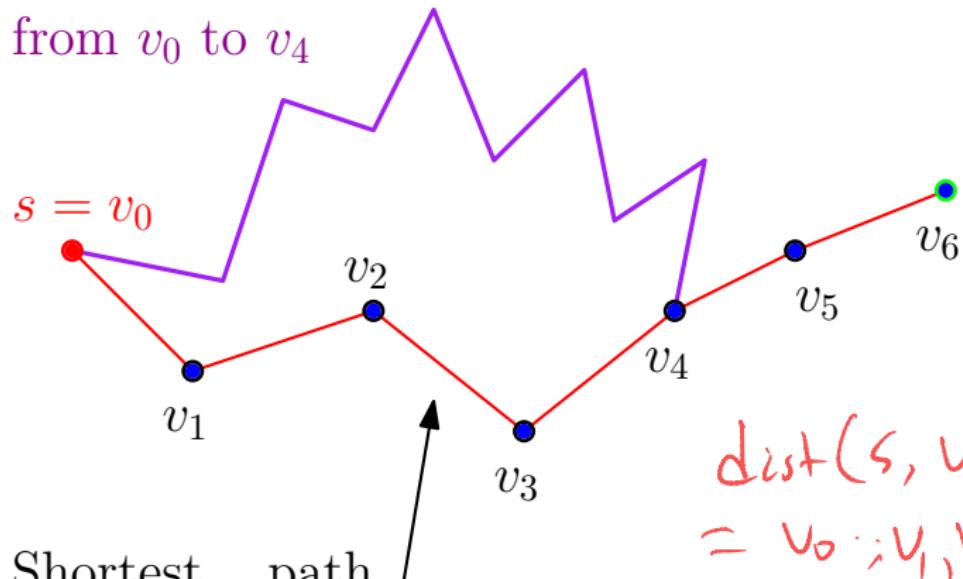
A proof by picture





A proof by picture

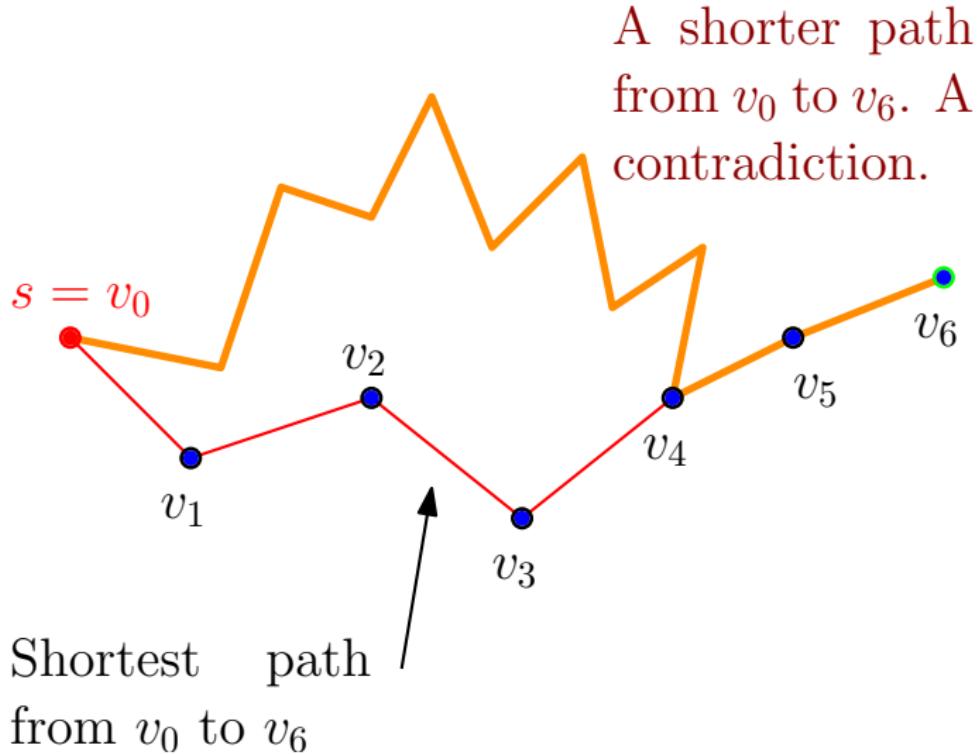
Shorter path
from v_0 to v_4



Shortest path
from v_0 to v_6

$$\begin{aligned} \text{dist}(s, v_4) \\ = \underline{\underline{v_0; v_1, v_2, v_3, v_4}} \end{aligned}$$

A proof by picture



A Basic Strategy

Explore vertices in increasing order of distance from s :

(For simplicity assume that nodes are at different distances from s and that no edge has zero length)

Initialize for each node v , $\text{dist}(s, v) = \infty$

Initialize $X = \{s\}$,

for $i = 2$ to $|V|$ do

(* Invariant: X contains the $i - 1$ closest nodes to s *)

Among nodes in $V - X$, find the node v that is the
 i 'th closest to s

Update $\text{dist}(s, v)$

$X = X \cup \{v\}$

A Basic Strategy

Explore vertices in increasing order of distance from s :

(For simplicity assume that nodes are at different distances from s and that no edge has zero length)

Initialize for each node v , $\text{dist}(s, v) = \infty$

Initialize $X = \{s\}$,

for $i = 2$ to $|V|$ do

(* Invariant: X contains the $i - 1$ closest nodes to s *)

Among nodes in $V - X$, find the node v that is the
 i 'th closest to s

Update $\text{dist}(s, v)$

$X = X \cup \{v\}$

How can we implement the step in the for loop?

✓

Finding the i th closest node

- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.

What do we know about the i th closest node?

Finding the i th closest node

- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.

What do we know about the i th closest node?

Claim

Let P be a shortest path from s to v where v is the i th closest node. Then, all intermediate nodes in P belong to X .

Finding the i th closest node

- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.

What do we know about the i th closest node?

Claim

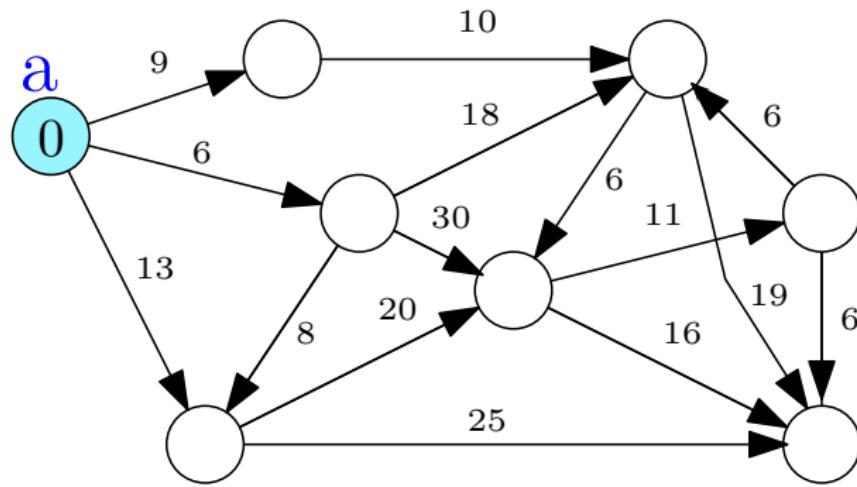
Let P be a shortest path from s to v where v is the i th closest node. Then, all intermediate nodes in P belong to X .

Proof.

If P had an intermediate node u not in X then u will be closer to s than v . Implies v is not the i 'th closest node to s - recall that X already has the $i - 1$ closest nodes. □

Finding the i th closest node repeatedly

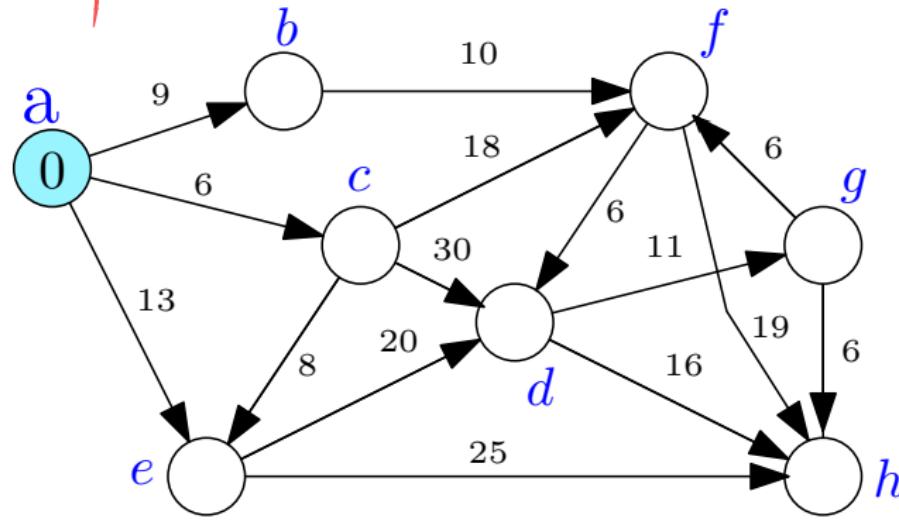
An example



Finding the i th closest node repeatedly

An example

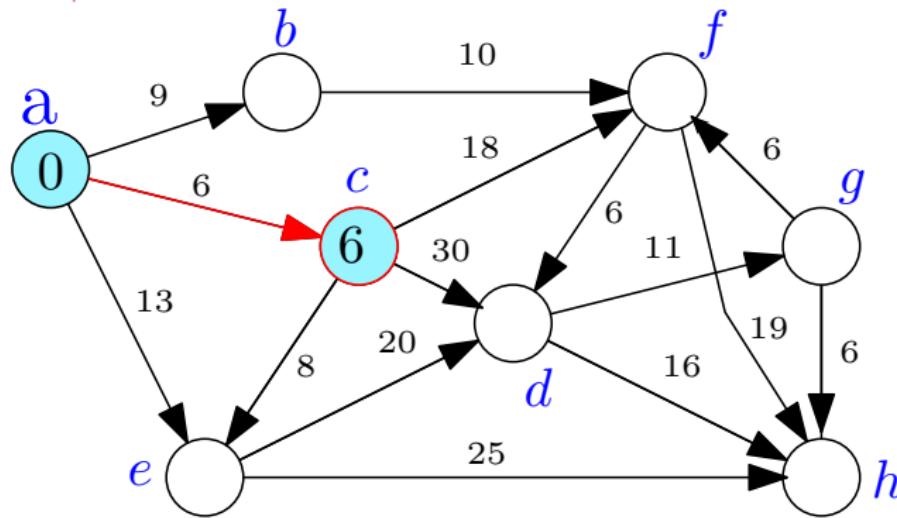
$$X = \{a\}$$



Finding the i th closest node repeatedly

An example

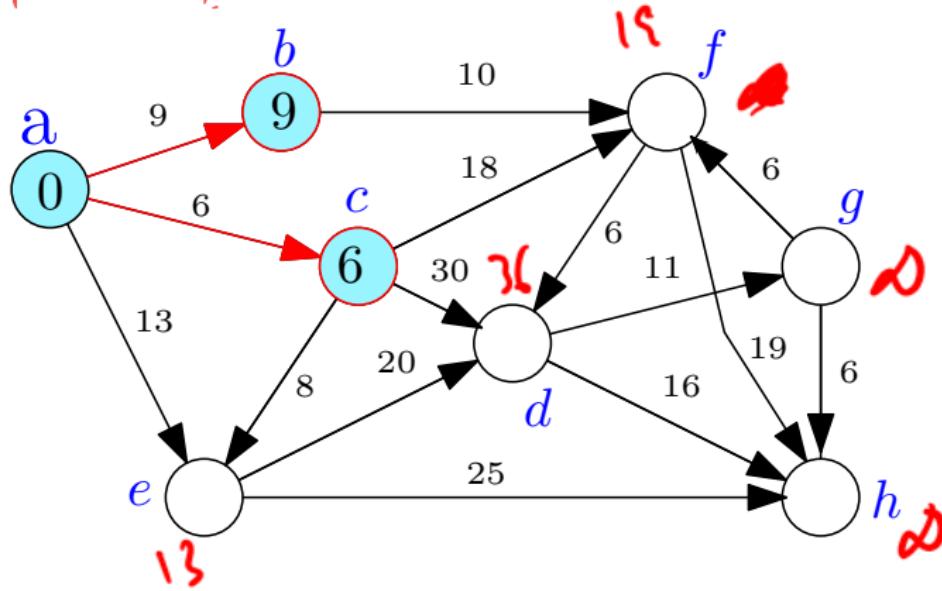
$$X = \{a, c\}$$



Finding the i th closest node repeatedly

An example

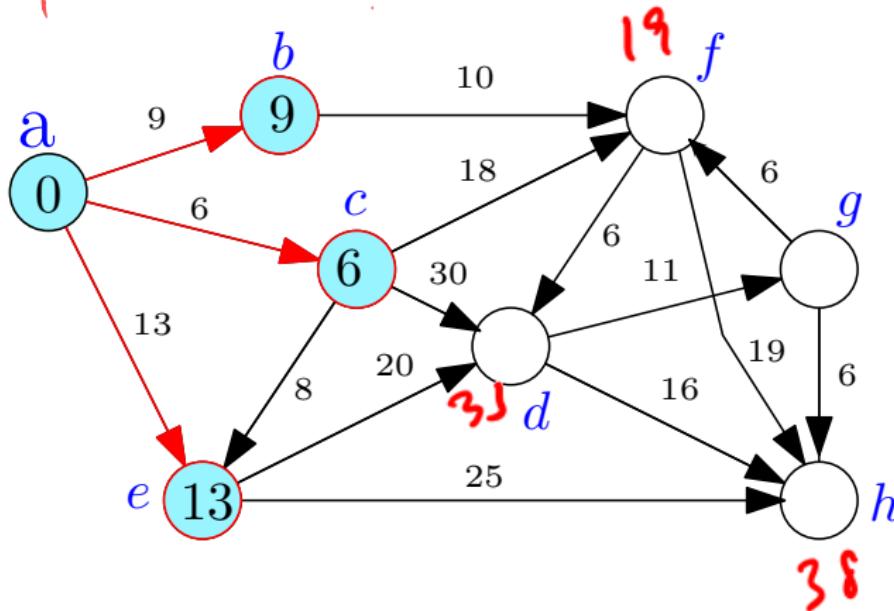
$$X = \{a, c, b\}$$



Finding the i th closest node repeatedly

An example

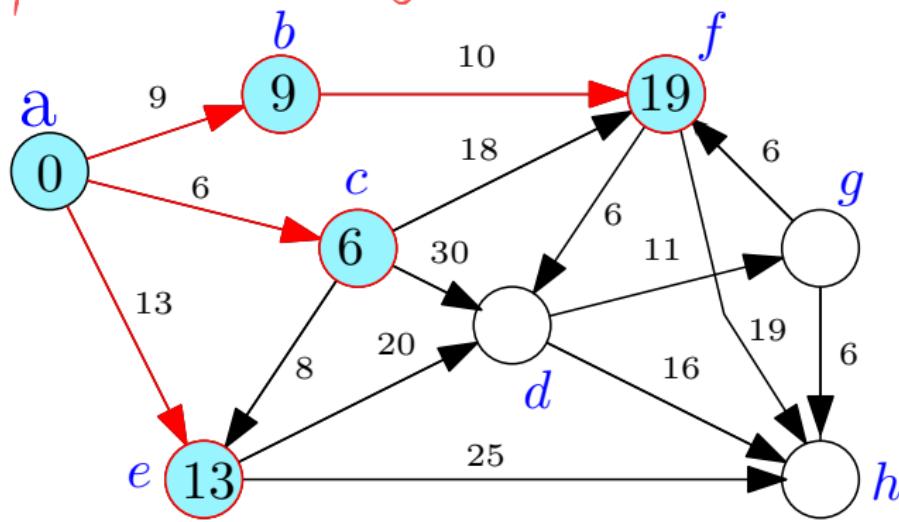
$$X = \{a, c, b, e\}$$



Finding the i th closest node repeatedly

An example

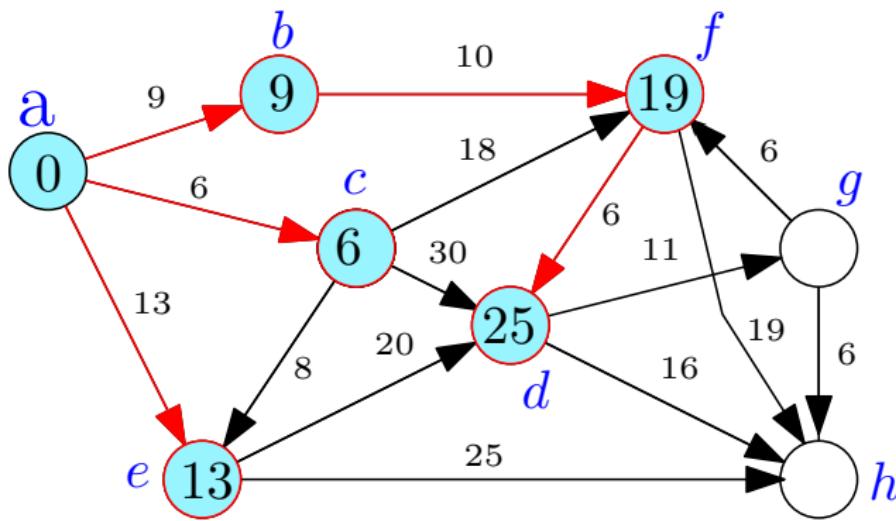
$$X = \{a, c, b, e, f\}$$



Finding the i th closest node repeatedly

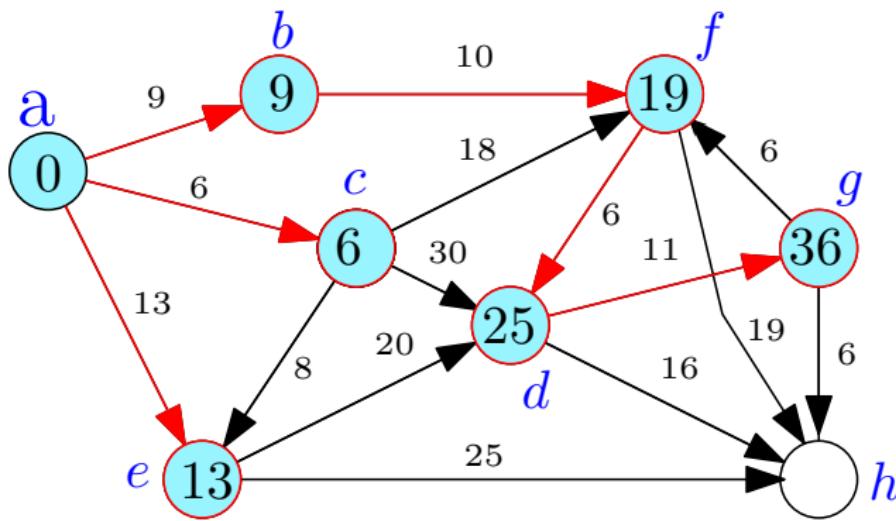
An example

$$X = \{a, c, b, e, f, d\}$$



Finding the i th closest node repeatedly

An example

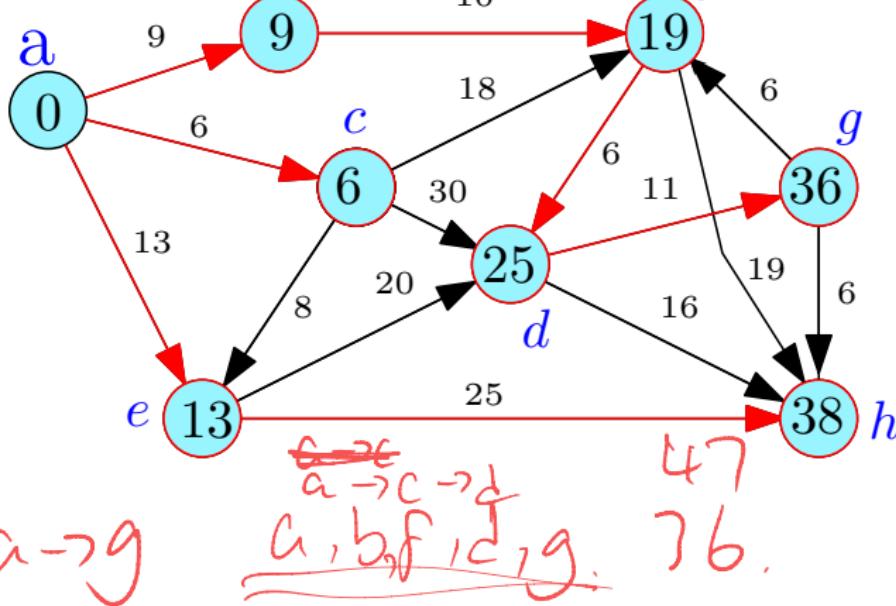


Finding the i th closest node repeatedly

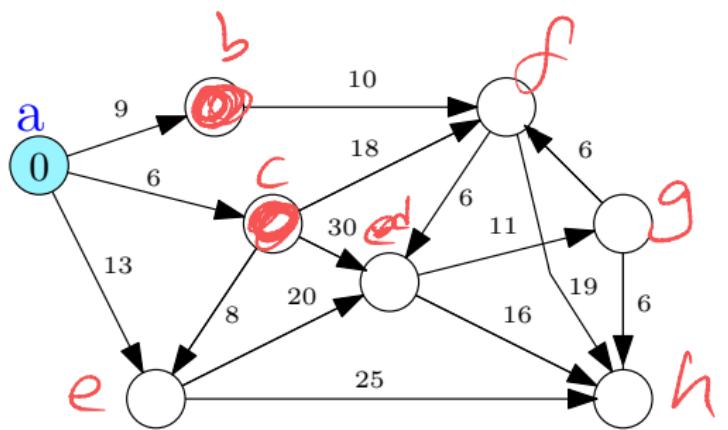
An example

$a \rightarrow d$ $\{a, b, f, d\}$, 25.

$b \rightarrow a, c, d, f$, 30.



Finding the i th closest node



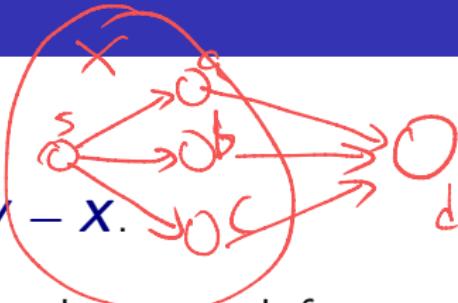
Corollary

The i th closest node is adjacent to X .

Finding the i th closest node

- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.
 $\textcolor{red}{S \rightarrow u}$.
- ③ For each $u \in V - X$ let $P(s, u, X)$ be a shortest path from s to u using only nodes in X as intermediate vertices.
- ④ Let $\underline{d'(s, u)}$ be the length of $P(s, u, X)$

Finding the i th closest node



- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.
- ① For each $u \in V - X$ let $P(s, u, X)$ be a shortest path from s to u using only nodes in X as intermediate vertices.
- ② Let $d'(s, u)$ be the length of $P(s, u, X)$

Observations: for each $u \in V - X$,

- ① $\text{dist}(s, u) \leq d'(s, u)$ since we are constraining the paths
- ② $d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$ - Why?

$$\min(\text{dist}(s \rightarrow t) + (t \rightarrow u))$$

Finding the i th closest node

- ① X contains the $i - 1$ closest nodes to s
- ② Want to find the i th closest node from $V - X$.
- ③ For each $u \in V - X$ let $P(s, u, X)$ be a shortest path from s to u using only nodes in X as intermediate vertices.
- ④ Let $d'(s, u)$ be the length of $P(s, u, X)$

Observations: for each $u \in V - X$,

- ① $\text{dist}(s, u) \leq d'(s, u)$ since we are constraining the paths
- ② $d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$ - Why?

Lemma

If v is the i th closest node to s , then $d'(s, v) = \text{dist}(s, v)$.

Finding the i th closest node

Lemma

Given:

- ① \mathbf{X} : Set of $i - 1$ closest nodes to s .
- ② $d'(s, u) = \min_{t \in \mathbf{X}} (\text{dist}(s, t) + \ell(t, u))$

If v is an i th closest node to s , then $d'(s, v) = \text{dist}(s, v)$.

Proof.

Let v be the i th closest node to s . Then there is a shortest path P from s to v that contains only nodes in \mathbf{X} as intermediate nodes (see previous claim). Therefore $d'(s, v) = \text{dist}(s, v)$. □

Finding the i th closest node

Lemma

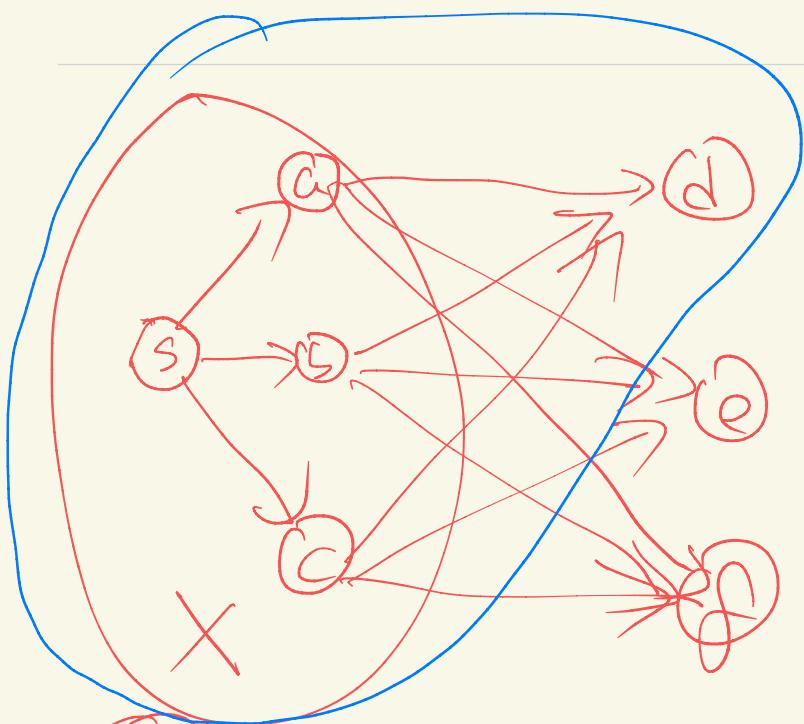
If v is an i th closest node to s , then $d'(s, v) = \text{dist}(s, v)$.

Corollary

The i th closest node to s is the node $v \in V - X$ such that $d'(s, v) = \min_{u \in V - X} d'(s, u)$.

Proof.

For every node $u \in V - X$, $\text{dist}(s, u) \leq d'(s, u)$ and for the i th closest node v , $\text{dist}(s, v) = d'(s, v)$. Moreover, $\text{dist}(s, u) \geq \text{dist}(s, v)$ for each $u \in V - X$. □



(i-1) 3个

最靠近S

的顶点

Algorithm

Initialize for each node v : $\text{dist}(s, v) = \infty$

Initialize $X = \emptyset$, $d'(s, s) = 0$

for $i = 1$ to $|V|$ **do**

(* Invariant: X contains the $i - 1$ closest nodes to s *)

(* Invariant: $d'(s, u)$ is shortest path distance from u to s using only X as intermediate nodes*)

Let v be such that $d'(s, v) = \min_{u \in V - X} d'(s, u)$

$\text{dist}(s, v) = d'(s, v)$

$X = X \cup \{v\}$

for each node u in $V - X$ **do**

$d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$

$$\min_{t \in V - X} d'(s, u) \rightarrow \underbrace{\text{dist}(s, t) + \text{dist}(t, u)}$$

Algorithm

```
Initialize for each node  $v$ :  $\text{dist}(s, v) = \infty$ 
Initialize  $X = \emptyset$ ,  $d'(s, s) = 0$ 
for  $i = 1$  to  $|V|$  do
    (* Invariant:  $X$  contains the  $i - 1$  closest nodes to  $s$  *)
    (* Invariant:  $d'(s, u)$  is shortest path distance from  $u$  to  $s$ 
       using only  $X$  as intermediate nodes*)
    Let  $v$  be such that  $d'(s, v) = \min_{u \in V - X} d'(s, u)$ 
     $\text{dist}(s, v) = d'(s, v)$ 
     $X = X \cup \{v\}$ 
    for each node  $u$  in  $V - X$  do
         $d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$ 
```

Correctness: By induction on i using previous lemmas.

Algorithm

```
Initialize for each node  $v$ :  $\text{dist}(s, v) = \infty$ 
Initialize  $X = \emptyset$ ,  $d'(s, s) = 0$ 
for  $i = 1$  to  $|V|$  do
    (* Invariant:  $X$  contains the  $i - 1$  closest nodes to  $s$  *)
    (* Invariant:  $d'(s, u)$  is shortest path distance from  $u$  to  $s$ 
       using only  $X$  as intermediate nodes*)
    Let  $v$  be such that  $d'(s, v) = \min_{u \in V - X} d'(s, u)$ 
     $\text{dist}(s, v) = d'(s, v)$ 
     $X = X \cup \{v\}$ 
    for each node  $u$  in  $V - X$  do
         $d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$ 
```

Correctness: By induction on i using previous lemmas.

Running time:

Algorithm

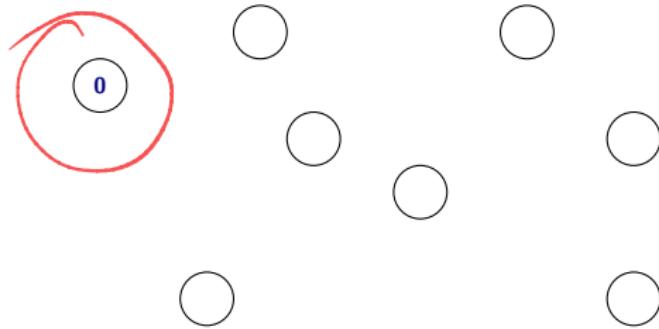
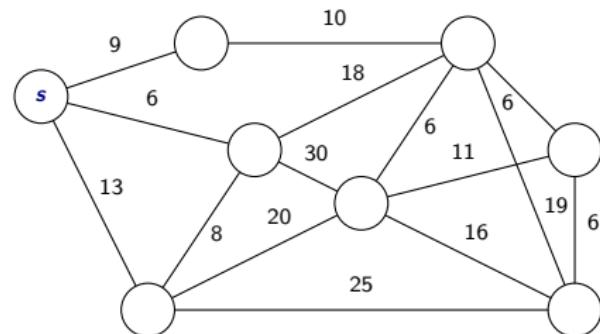
```
Initialize for each node  $v$ :  $\text{dist}(s, v) = \infty$ 
Initialize  $X = \emptyset$ ,  $d'(s, s) = 0$ 
for  $i = 1$  to  $|V|$  do
    (* Invariant:  $X$  contains the  $i - 1$  closest nodes to  $s$  *)
    (* Invariant:  $d'(s, u)$  is shortest path distance from  $u$  to  $s$ 
       using only  $X$  as intermediate nodes*)
    Let  $v$  be such that  $d'(s, v) = \min_{u \in V - X} d'(s, u)$ 
     $\text{dist}(s, v) = d'(s, v)$ 
     $X = X \cup \{v\}$ 
    for each node  $u$  in  $V - X$  do
         $d'(s, u) = \min_{t \in X} (\text{dist}(s, t) + \ell(t, u))$ 
```

Correctness: By induction on i using previous lemmas.

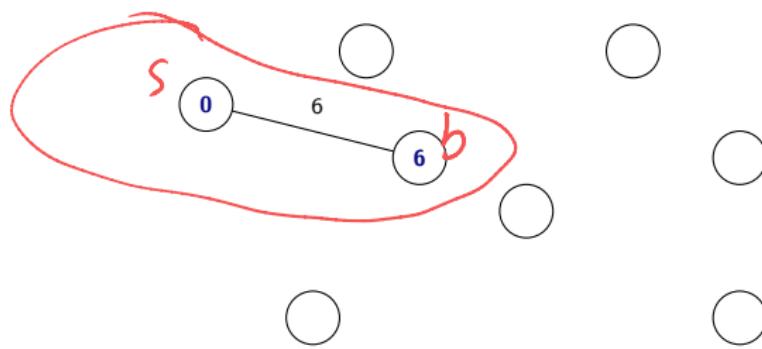
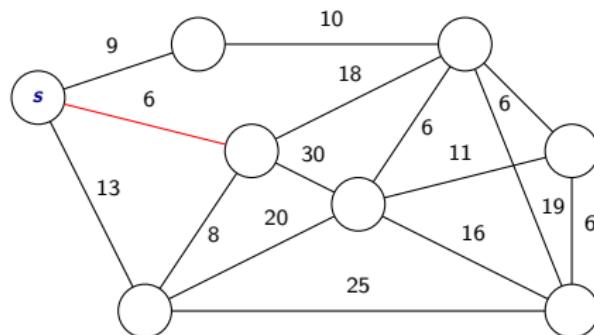
Running time: $O(n \cdot (n + m))$ time.

- ➊ n outer iterations. In each iteration, $d'(s, u)$ for each u by scanning all edges out of nodes in X ; $O(m + n)$ time/iteration.

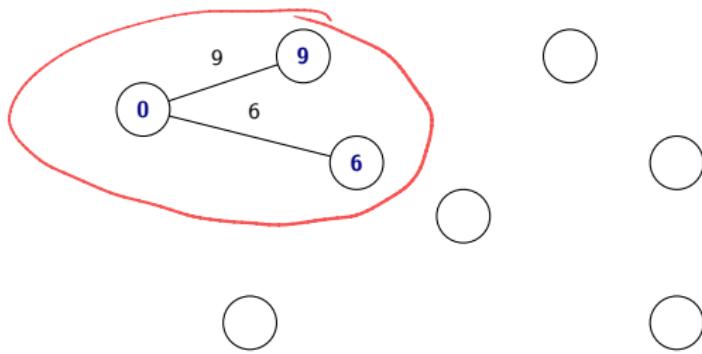
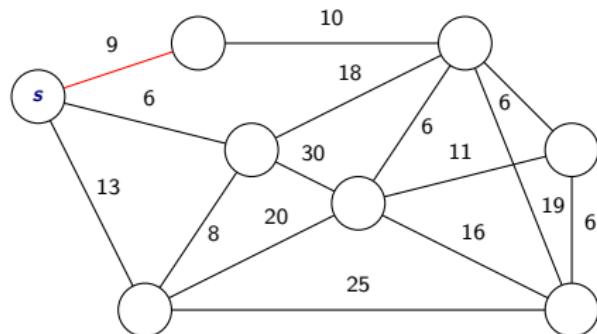
Example



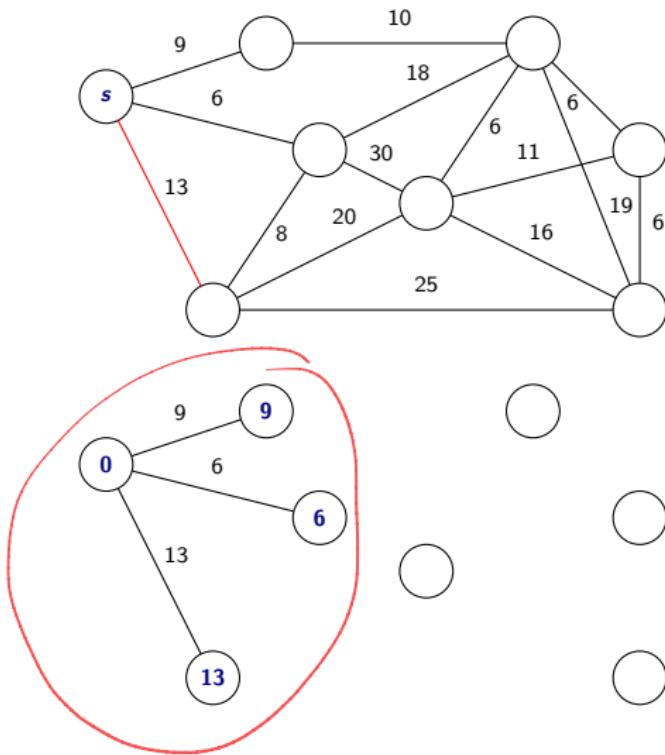
Example



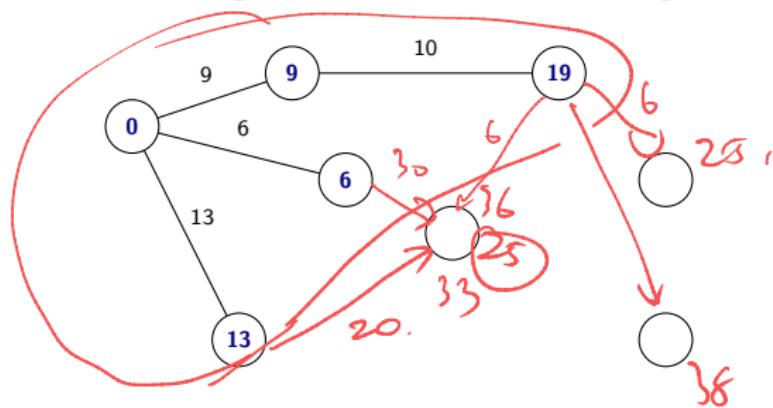
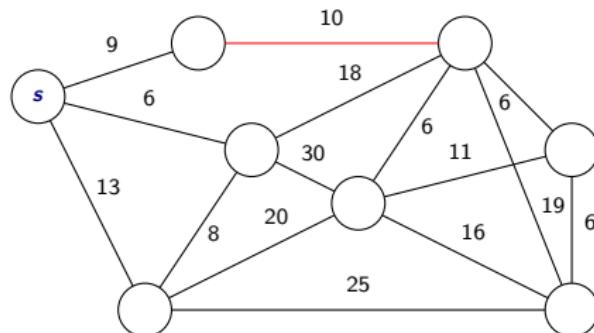
Example



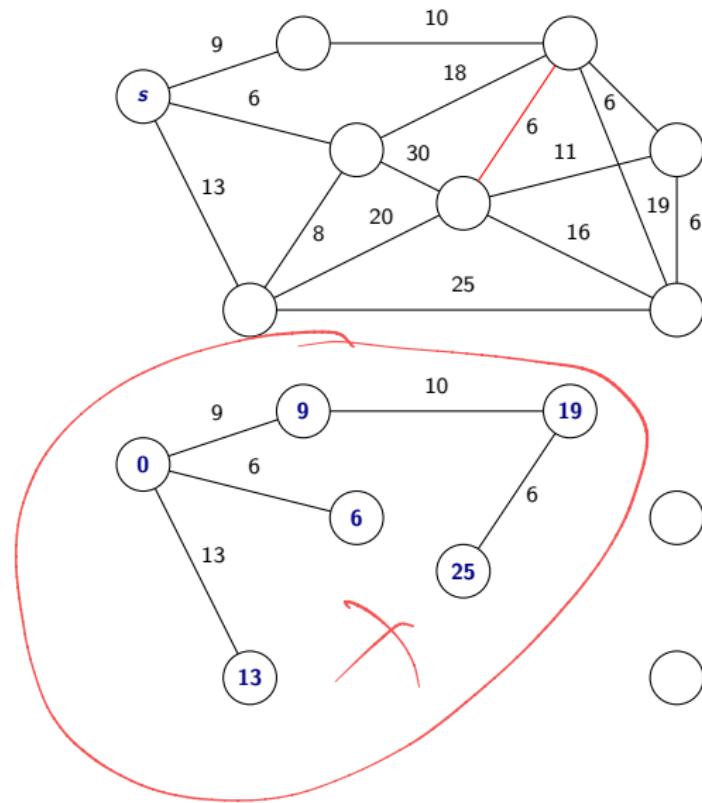
Example



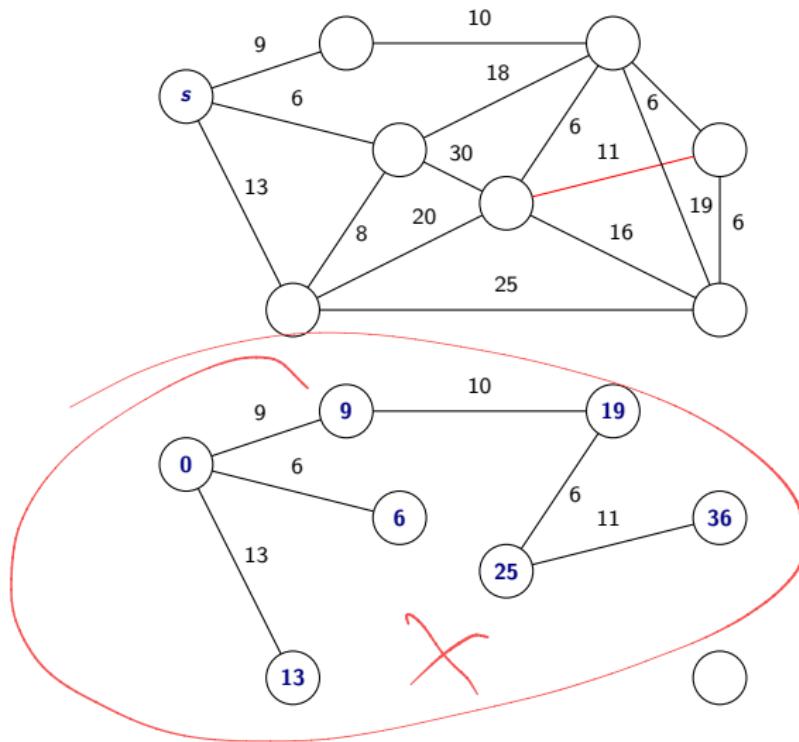
Example



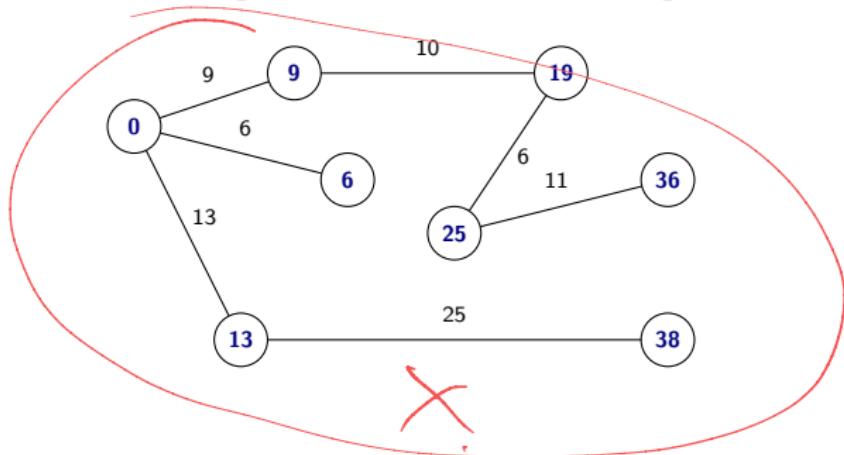
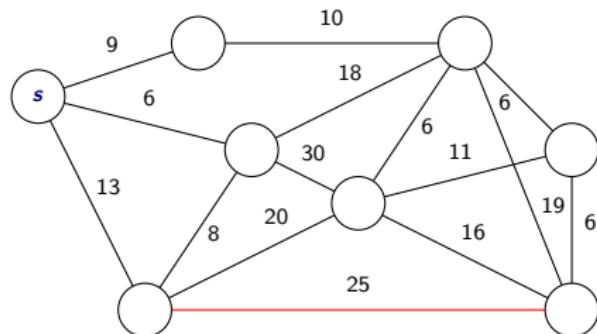
Example



Example



Example



Improved Algorithm

- ① Main work is to compute the $d'(s, u)$ values in each iteration
- ② $d'(s, u)$ changes from iteration i to $i + 1$ only because of the node v that is added to X in iteration i .

Improved Algorithm

- ① Main work is to compute the $d'(s, u)$ values in each iteration
- ② $d'(s, u)$ changes from iteration i to $i + 1$ only because of the node v that is added to X in iteration i .

```
Initialize for each node  $v$ ,  $\text{dist}(s, v) = d'(s, v) = \infty$ 
```

```
Initialize  $X = \emptyset$ ,  $d'(s, s) = 0$ 
```

```
for  $i = 1$  to  $|V|$  do
```

// X contains the $i - 1$ closest nodes to s ,

// and the values of $d'(s, u)$ are current

```
Let  $v$  be node realizing  $d'(s, v) = \min_{u \in V - X} d'(s, u)$ 
```

```
 $\text{dist}(s, v) = d'(s, v)$ 
```

```
 $X = X \cup \{v\}$ 
```

```
Update  $d'(s, u)$  for each  $u$  in  $V - X$  as follows:
```

$$d'(s, u) = \min(d'(s, u), \text{dist}(s, v) + \ell(v, u))$$

Running time:

Improved Algorithm

Initialize for each node v , $\text{dist}(s, v) = d'(s, v) = \infty$

Initialize $X = \emptyset$, $d'(s, s) = 0$

for $i = 1$ to $|V|$ **do**

- // X contains the $i - 1$ closest nodes to s ,
- // and the values of $d'(s, u)$ are current

Let v be node realizing $d'(s, v) = \min_{u \in V - X} d'(s, u)$

$\text{dist}(s, v) = d'(s, v)$

$X = X \cup \{v\}$

Update $d'(s, u)$ for each u in $V - X$ as follows:

$$d'(s, u) = \min(d'(s, u), \text{dist}(s, v) + \ell(v, u))$$

Running time: $O(m + n^2)$ time.

- ① n outer iterations and in each iteration following steps
- ② updating $d'(s, u)$ after v is added takes $O(\deg(v))$ time so total work is $O(m)$ since a node enters X only once
- ③ Finding v from $d'(s, u)$ values is $O(n)$ time

Dijkstra's Algorithm

- ➊ eliminate $d'(s, u)$ and let $\text{dist}(s, u)$ maintain it
- ➋ update dist values after adding v by scanning edges out of v

```
1 Initialize for each node  $v$ ,  $\text{dist}(s, v) = \infty$ 
2 Initialize  $X = \emptyset$ ,  $\text{dist}(s, s) = 0$ 
3 for  $i = 1$  to  $|V|$  do
    Let  $v$  be such that  $\text{dist}(s, v) = \min_{u \in V - X} \text{dist}(s, u)$ 
     $X = X \cup \{v\}$ 
    for each  $u$  in  $\text{Adj}(v)$  do
         $\text{dist}(s, u) = \min(\text{dist}(s, u), \text{dist}(s, v) + \ell(v, u))$ 
```

Priority Queues to maintain dist values for faster running time

Dijkstra's Algorithm

用优先队列: $n \cdot (\log n + m \cdot \log n)$
 $= n \log n + m \log n$

- ① eliminate $d'(s, u)$ and let $\text{dist}(s, u)$ maintain it

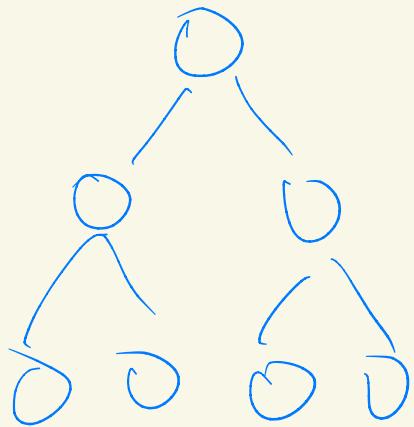
- ② update dist values after adding v by scanning edges out of v

1 Initialize for each node v , $\text{dist}(s, v) = \infty$
2 Initialize $X = \emptyset$, $\text{dist}(s, s) = 0$ 不用的: $O(n \cdot (n + m_n))$
for $i = 1$ to $|V|$ do
 3 $X = X \cup \{v\}$ $= O(n^2)$
 4 Let v be such that $\text{dist}(s, v) = \min_{u \in V - X} \text{dist}(s, u)$
 5 $X = X \cup \{v\}$ $= O(n^2)$
 6 for each u in $\text{Adj}(v)$ do
 7 $\text{dist}(s, u) = \min(\text{dist}(s, u), \text{dist}(s, v) + \ell(v, u))$ $= O(n)$

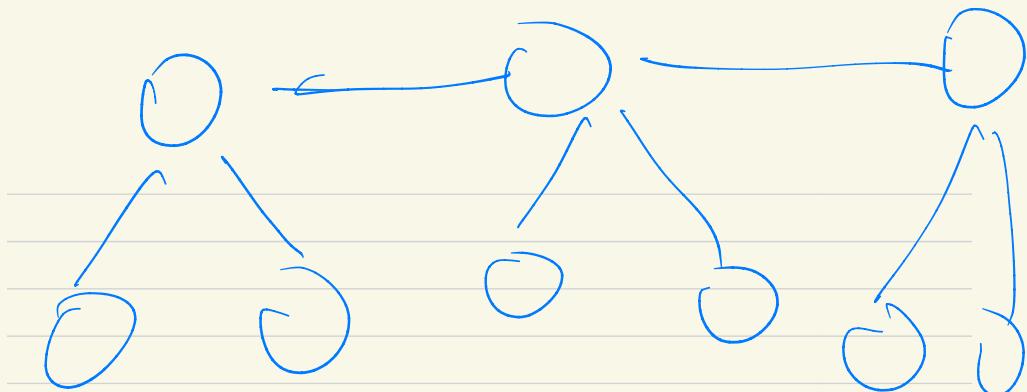
Priority Queues to maintain dist values for faster running time

- ① Using heaps and standard priority queues: $O((m + n) \log n)$
- ② Using Fibonacci heaps: $O(m + n \log n)$.

$m \log n + n \log n$



$\log n$.



Fibonacci

Heap

不用优先队列

3, $O(n \cdot k)$

4, $O(n)$

5, $O(\frac{m}{n} \cdot n)$

6, $O(1)$

$$O(n \cdot (n + m/n - 1))$$

$$= O(n^2 + m)$$

$$\approx O(n^2)$$

用优先队列 (二叉堆)

3, $O(n)$

4, $O(n \log m)$

$$O(n \cdot (\log n +$$

5, $O(\frac{m}{n})$

$$m/n \cdot \log n))$$

6, $O(\log m)$

$$= O(n \log n + m \log m)$$

[#] Fibonacci heap:

3, $O(N)$

4, $O(\log N)$

6, $O(m/n)$

7, $O(1)$

$O(N \cdot (\log N + m/n))$

$= O(N \log N + m)$

Priority Queues

Data structure to store a set S of n elements where each element $v \in S$ has an associated real/integer key $k(v)$ such that the following operations:

- ① **makePQ**: create an empty queue.
- ② **findMin**: find the minimum key in S .
- ③ **extractMin**: Remove $v \in S$ with smallest key and return it.
- ④ **insert($v, k(v)$)**: Add new element v with key $k(v)$ to S .
- ⑤ **delete(v)**: Remove element v from S .

Priority Queues

Data structure to store a set S of n elements where each element $v \in S$ has an associated real/integer key $k(v)$ such that the following operations:

- ① **makePQ**: create an empty queue.
- ② **findMin**: find the minimum key in S .
- ③ **extractMin**: Remove $v \in S$ with smallest key and return it.
- ④ **insert($v, k(v)$)**: Add new element v with key $k(v)$ to S .
- ⑤ **delete(v)**: Remove element v from S .
- ⑥ **decreaseKey($v, k'(v)$)**: decrease key of v from $k(v)$ (current key) to $k'(v)$ (new key). Assumption: $k'(v) \leq k(v)$.
- ⑦ **meld**: merge two separate priority queues into one.

Priority Queues

Data structure to store a set S of n elements where each element $v \in S$ has an associated real/integer key $k(v)$ such that the following operations:

- ① **makePQ**: create an empty queue.
- ② **findMin**: find the minimum key in S .
- ③ **extractMin**: Remove $v \in S$ with smallest key and return it.
- ④ **insert($v, k(v)$)**: Add new element v with key $k(v)$ to S .
- ⑤ **delete(v)**: Remove element v from S .
- ⑥ **decreaseKey($v, k'(v)$)**: decrease key of v from $k(v)$ (current key) to $k'(v)$ (new key). Assumption: $k'(v) \leq k(v)$.
- ⑦ **meld**: merge two separate priority queues into one.

All operations can be performed in $O(\log n)$ time.

decreaseKey is implemented via **delete** and **insert**.

Dijkstra's Algorithm using Priority Queues

```
Q ← makePQ()
insert(Q, (s, 0))
for each node u ≠ s do
    insert(Q, (u, ∞))
X ← ∅
for i = 1 to |V| do
    (v, dist(s, v)) = extractMin(Q)
    X = X ∪ {v}
    for each u in Adj(v) do
        decreaseKey(Q, (u, min(dist(s, u), dist(s, v) + ℓ(v, u)))) .
```

Priority Queue operations:

- ① $O(n)$ **insert** operations
- ② $O(n)$ **extractMin** operations
- ③ $O(m)$ **decreaseKey** operations

Implementing Priority Queues via Heaps

Using Heaps

Store elements in a heap based on the key value

- ➊ All operations can be done in $O(\log n)$ time

Implementing Priority Queues via Heaps

Using Heaps

Store elements in a heap based on the key value

- ➊ All operations can be done in $O(\log n)$ time

Dijkstra's algorithm can be implemented in $O((n + m) \log n)$ time.

Priority Queues: Fibonacci Heaps/Relaxed Heaps

Fibonacci Heaps

- ① **extractMin, insert, delete, meld** in $O(\log n)$ time
- ② **decreaseKey** in $O(1)$ *amortized* time:

Priority Queues: Fibonacci Heaps/Relaxed Heaps

Fibonacci Heaps

- ① **extractMin, insert, delete, meld** in $O(\log n)$ time
- ② **decreaseKey** in $O(1)$ amortized time: ℓ **decreaseKey** operations for $\ell \geq n$ take together $O(\ell)$ time
- ③ Relaxed Heaps: **decreaseKey** in $O(1)$ worst case time but at the expense of **meld** (not necessary for Dijkstra's algorithm)

Priority Queues: Fibonacci Heaps/Relaxed Heaps

Fibonacci Heaps

- ① **extractMin, insert, delete, meld** in $O(\log n)$ time
 - ② **decreaseKey** in $O(1)$ amortized time: ℓ **decreaseKey** operations for $\ell \geq n$ take together $O(\ell)$ time
 - ③ Relaxed Heaps: **decreaseKey** in $O(1)$ worst case time but at the expense of **meld** (not necessary for Dijkstra's algorithm)
-
- ① Dijkstra's algorithm can be implemented in $O(n \log n + m)$ time. If $m = \Omega(n \log n)$, running time is linear in input size.

Priority Queues: Fibonacci Heaps/Relaxed Heaps

Fibonacci Heaps

- ① **extractMin, insert, delete, meld** in $O(\log n)$ time
- ② **decreaseKey** in $O(1)$ amortized time: ℓ **decreaseKey** operations for $\ell \geq n$ take together $O(\ell)$ time
- ③ Relaxed Heaps: **decreaseKey** in $O(1)$ worst case time but at the expense of **meld** (not necessary for Dijkstra's algorithm)

- ① Dijkstra's algorithm can be implemented in $O(n \log n + m)$ time. If $m = \Omega(n \log n)$, running time is linear in input size.
- ② Data structures are complicated to analyze/implement. Recent work has obtained data structures that are easier to analyze and implement, and perform well in practice. Rank-Pairing Heaps (European Symposium on Algorithms, September 2009!)

Shortest Path Tree

Dijkstra's algorithm finds the shortest path distances from s to V .

Question: How do we find the paths themselves?

Shortest Path Tree

Dijkstra's algorithm finds the shortest path distances from s to V .

Question: How do we find the paths themselves?

```
Q = makePQ()
insert(Q, (s, 0))
prev(s) ← null
for each node u ≠ s do
    insert(Q, (u, ∞))
    prev(u) ← null

X = ∅
for i = 1 to |V| do
    (v, dist(s, v)) = extractMin(Q)
    X = X ∪ {v}
    for each u in Adj(v) do
        if (dist(s, v) + ℓ(v, u) < dist(s, u)) then
            decreaseKey(Q, (u, dist(s, v) + ℓ(v, u)))
            prev(u) = v
```

Shortest Path Tree

Lemma

The edge set $(u, \text{prev}(u))$ is the reverse of a shortest path tree rooted at s . For each u , the reverse of the path from u to s in the tree is a shortest path from s to u .

Proof Sketch.

- ① The edge set $\{(u, \text{prev}(u)) \mid u \in V\}$ induces a directed in-tree rooted at s (Why?)
- ② Use induction on $|X|$ to argue that the tree is a shortest path tree for nodes in V .



Shortest paths to s

Dijkstra's algorithm gives shortest paths from s to all nodes in V .
How do we find shortest paths from all of V to s ?

Shortest paths to s

Dijkstra's algorithm gives shortest paths from s to all nodes in V .
How do we find shortest paths from all of V to s ?

- ① In undirected graphs shortest path from s to u is a shortest path from u to s so there is no need to distinguish.
- ② In directed graphs, use Dijkstra's algorithm in G^{rev} !

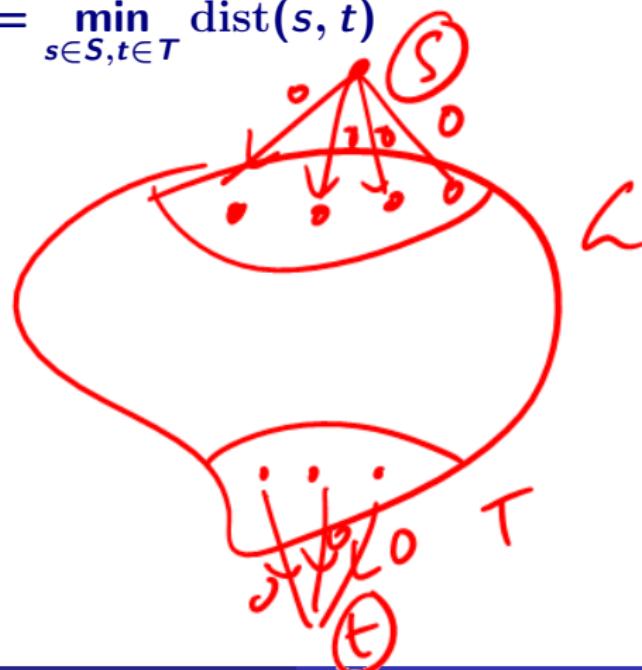
G^{rev} \Rightarrow transpose Graph.
把边的走向反转

Shortest paths between sets of nodes

Suppose we are given $S \subset V$ and $T \subset V$. Want to find shortest path from S to T defined as:

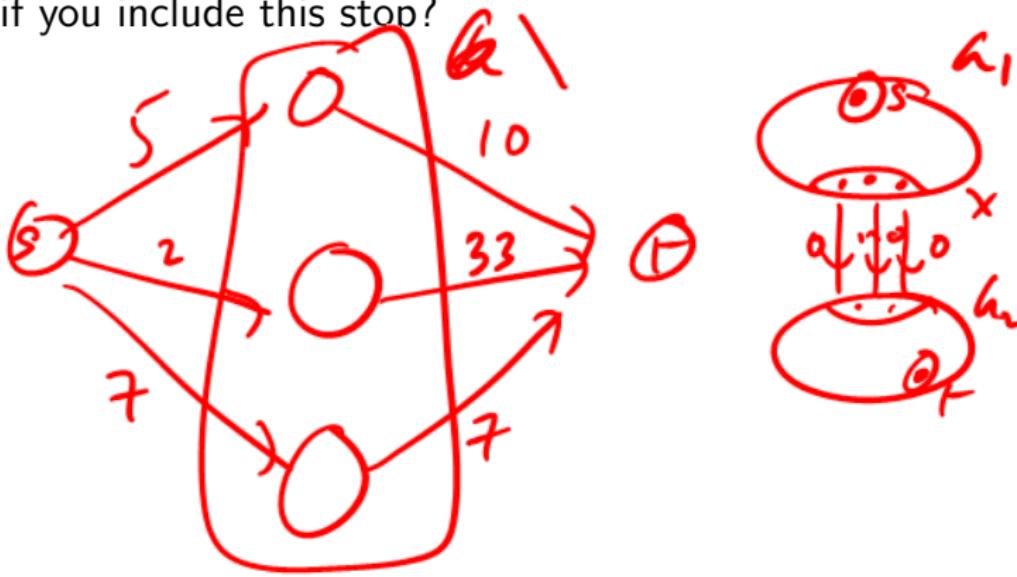
$$\text{dist}(S, T) = \min_{s \in S, t \in T} \text{dist}(s, t)$$

How do we find $\text{dist}(S, T)$?



Example Problem

You want to go from your house to a friend's house. Need to pick up some dessert along the way and hence need to stop at one of the many potential stores along the way. How do you calculate the "shortest" trip if you include this stop?



Example Problem

You want to go from your house to a friend's house. Need to pick up some dessert along the way and hence need to stop at one of the many potential stores along the way. How do you calculate the "shortest" trip if you include this stop?

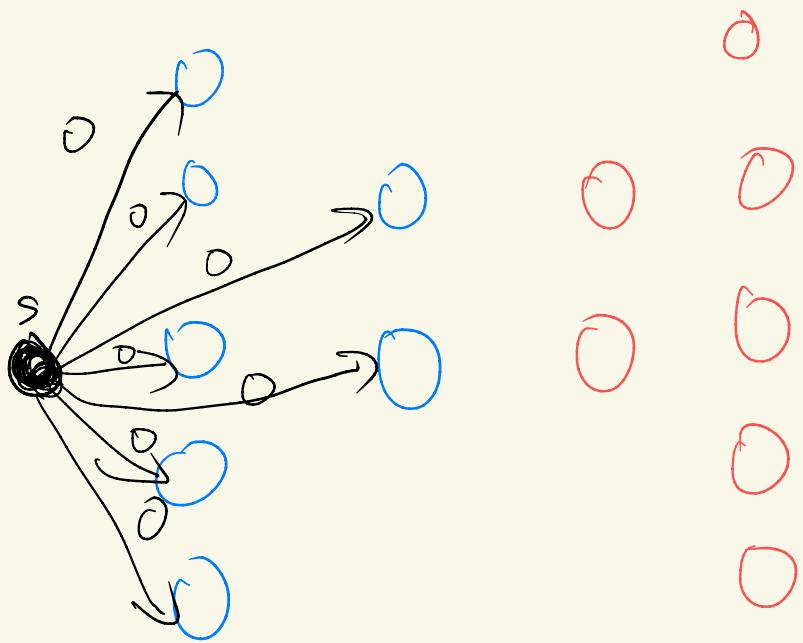
Given $G = (V, E)$ and edge lengths $\ell(e), e \in E$. Want to go from s to t . A subset $X \subset V$ that corresponds to stores. Want to find $\min_{x \in X} d(s, x) + d(x, t)$.

Example Problem

You want to go from your house to a friend's house. Need to pick up some dessert along the way and hence need to stop at one of the many potential stores along the way. How do you calculate the "shortest" trip if you include this stop?

Given $G = (V, E)$ and edge lengths $\ell(e), e \in E$. Want to go from s to t . A subset $X \subset V$ that corresponds to stores. Want to find $\min_{x \in X} d(s, x) + d(x, t)$.

Basic solution: Compute for each $x \in X$, $d(s, x)$ and $d(x, t)$ and take minimum. $2|X|$ shortest path computations.
 $O(|X|(m + n \log n))$.



Log N

Example Problem

You want to go from your house to a friend's house. Need to pick up some dessert along the way and hence need to stop at one of the many potential stores along the way. How do you calculate the "shortest" trip if you include this stop?

Given $G = (V, E)$ and edge lengths $\ell(e), e \in E$. Want to go from s to t . A subset $X \subset V$ that corresponds to stores. Want to find $\min_{x \in X} d(s, x) + d(x, t)$.

Basic solution: Compute for each $x \in X$, $d(s, x)$ and $d(x, t)$ and take minimum. $2|X|$ shortest path computations.
 $O(|X|(m + n \log n))$.

Better solution: Compute shortest path distances from s to every node $v \in V$ with one Dijkstra. Compute from every node $v \in V$ shortest path distance to t with one Dijkstra.